



Applied Machine Learning

Amin Nazari

Spring 2025

Chapter 1

Supervised Learning


Linear Regression

Outline

- Linear Regression
- Correlation matrix
- Loss function and Gradient descent
- Learning Rate
- Non-linear regression
- Overfitting
 - Train-test split
 - Regularization

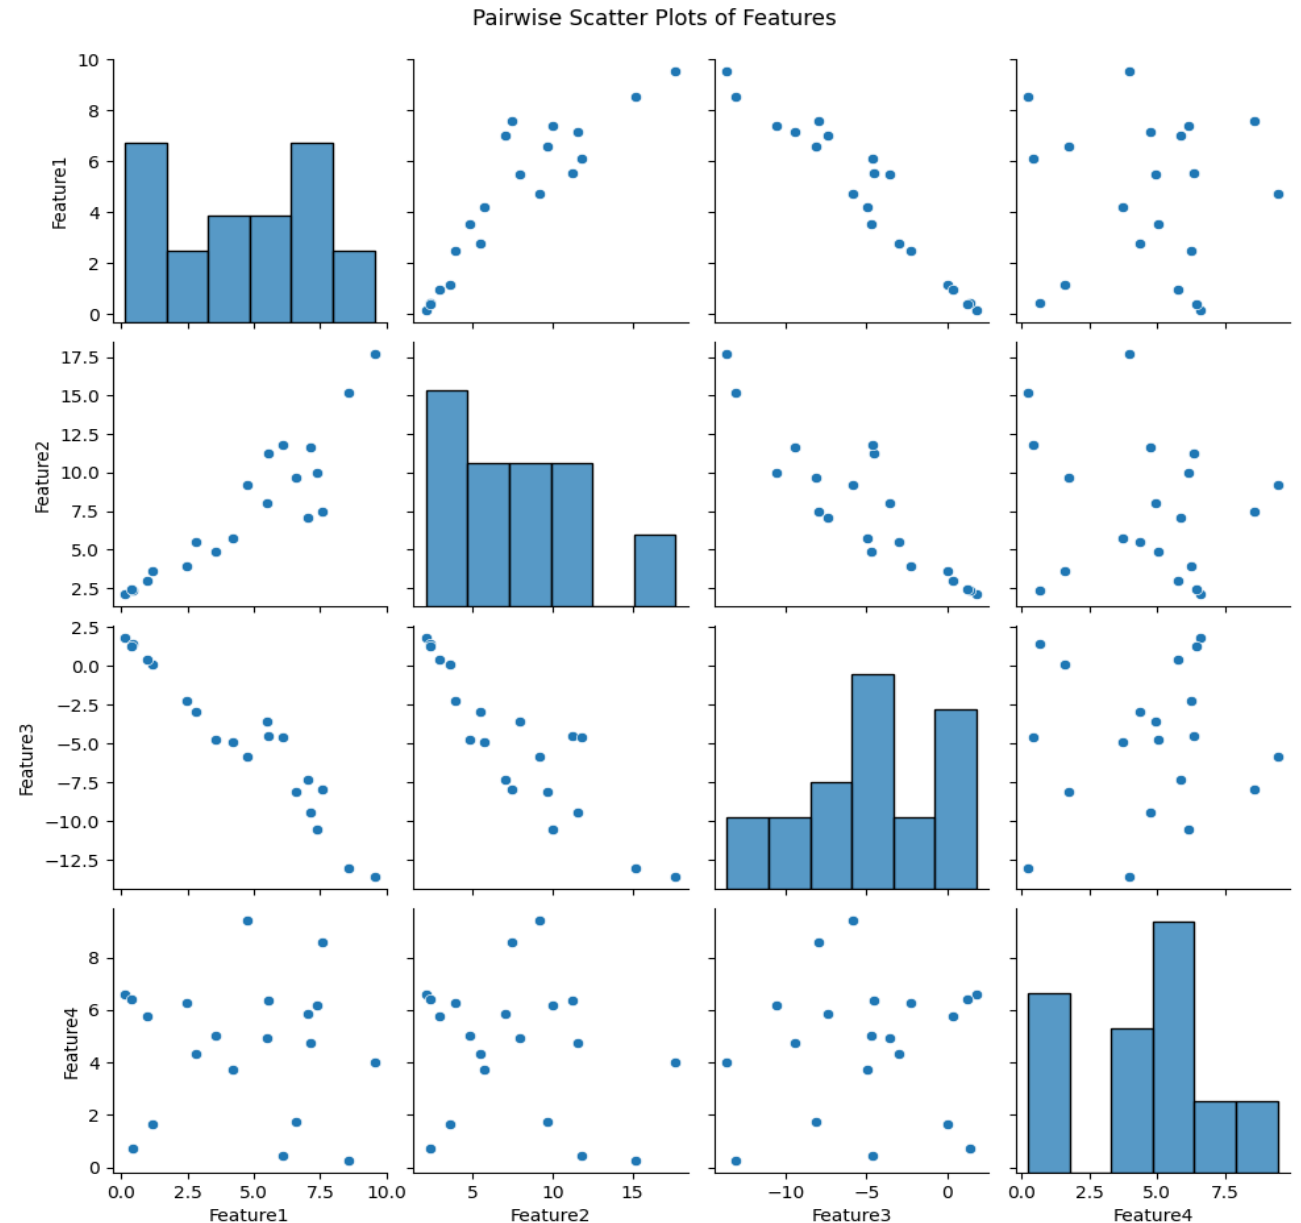
Linear Regression

- The goal of linear regression is to find a linear relationship between **one or more independent variables** and a **dependent variable**.
- A linear regression uses the fitting of a **line** or **hyperplane** to model the relationship between the variables.
- This linear model can then be used to predict the value of the dependent variable for new data.



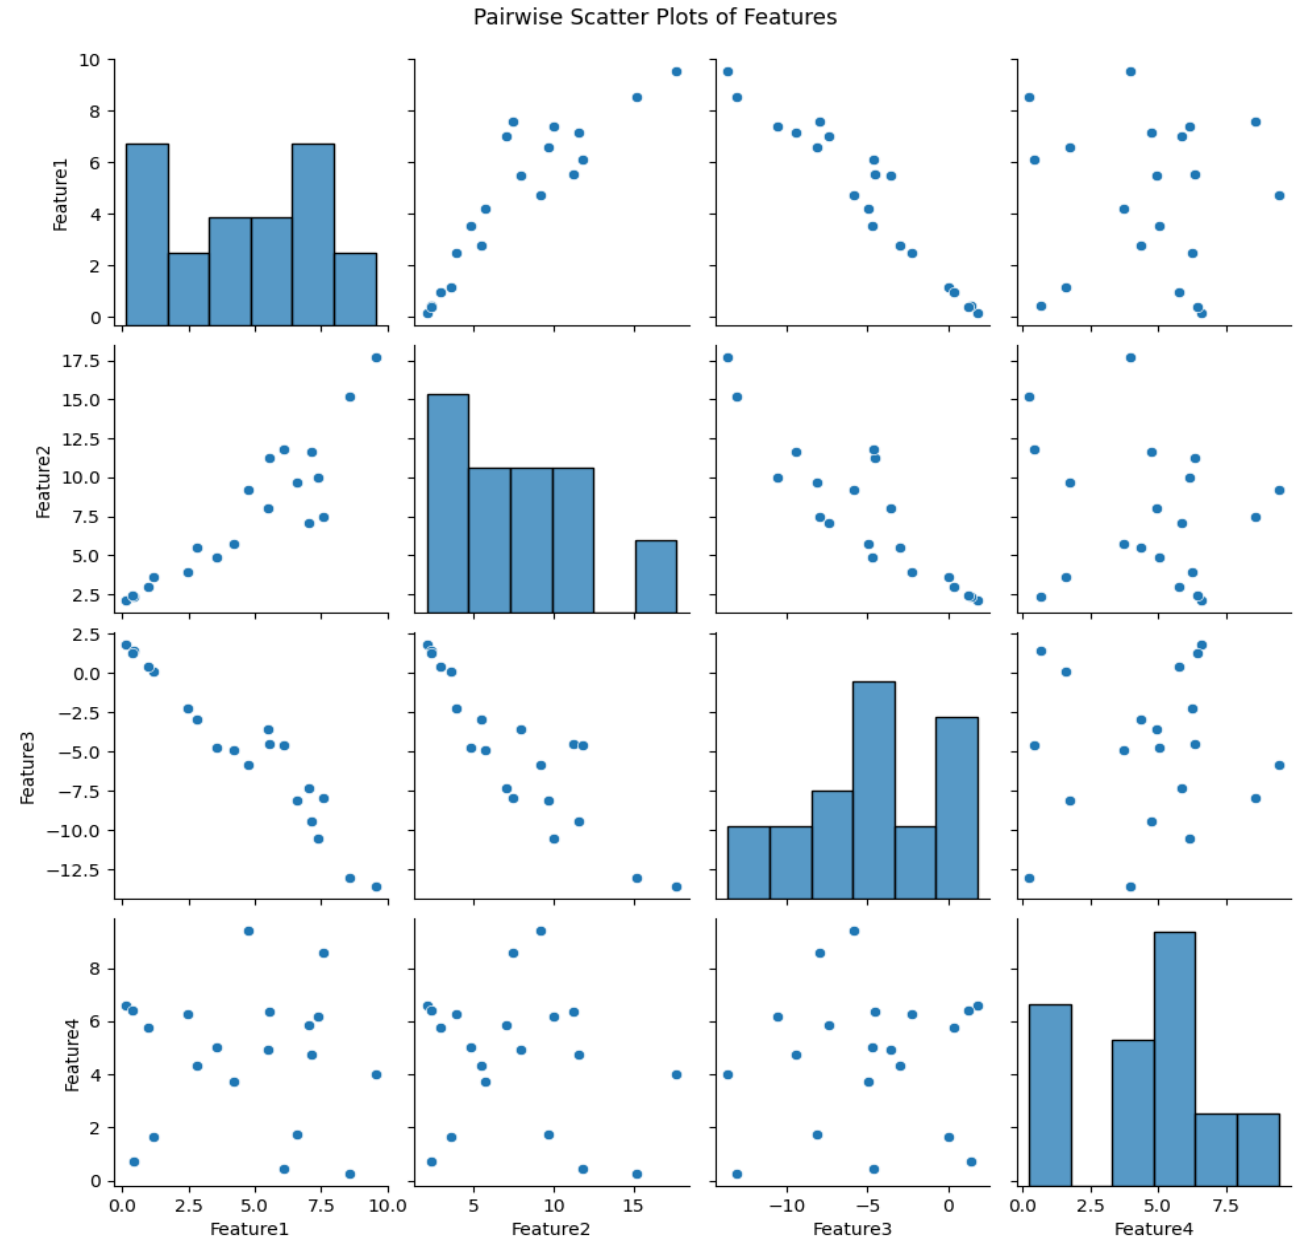

Linear Regression

- Is linear regression feasible?



Correlation

$$r = \frac{n\sum xy - (\sum x)(\sum y)}{\sqrt{n\sum x^2 - (\sum x^2)}\sqrt{n\sum y^2 - (\sum y^2)}}$$




Correlation (Statistical significance)

$$r = \frac{n\sum xy - (\sum x)(\sum y)}{\sqrt{n\sum x^2 - (\sum x)^2}\sqrt{n\sum y^2 - (\sum y)^2}}$$

H0: $\rho = 0$ (meaning no relationship)


H1: $\rho \neq 0$ (relationship exists)

If our test value (t) falls outside this range (-2.262, 2.262), we can reject our null hypothesis. (With 95% confidence)



$$t = \frac{r}{\sqrt{\frac{1-r^2}{n-2}}}$$

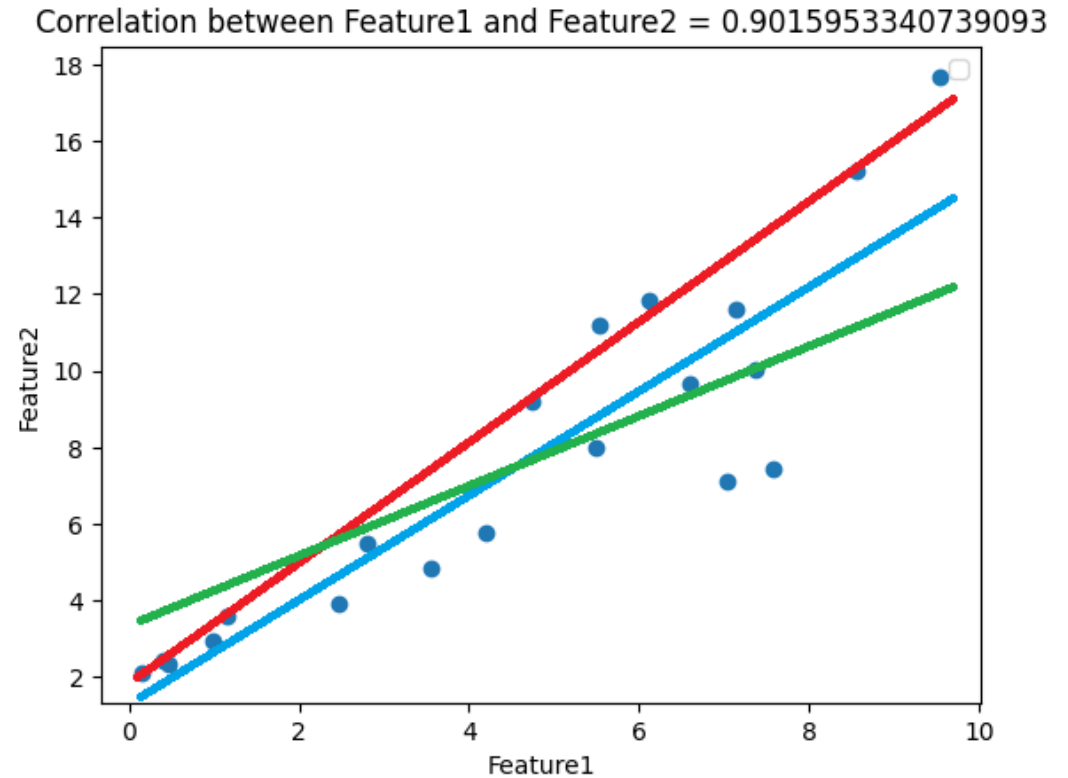

t-distribution and normal-distribution



Python Code

- Correlation Matrix
- Statistical significance
- p-value describes the likelihood of data occurring if the null hypothesis were true

Which line is best?




Model and loss function

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \cdots + \theta_n x_n$$

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

- m is the number of training examples.
- $y^{(i)}$ is the actual value for the i -th training example.
- $h_{\theta}(x^{(i)})$ is the predicted value for the i -th training example.




Gradient Descent

$\forall \theta_j = \text{Random initialize}$

$$\theta_j := \theta_j - \alpha \frac{\partial J(\theta)}{\partial \theta_j}$$

$$\frac{\partial J(\theta)}{\partial \theta_j} = \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$



Gradient Descent

Initialize model parameters (weights and biases) randomly

Set learning rate (α)

For each iteration:

- Initialize total_loss = 0

- Initialize gradients for weights and biases as zero

- For each data point (x_i, y_i) in the training set:

 - predicted_y_i = model(x_i , weights, biases)

 - gradient_weights, gradient_biases = compute_gradients($x_i, y_i, \text{predicted_y_i}$)

 - gradients_weights += gradient_weights

 - gradients_biases += gradient_biases

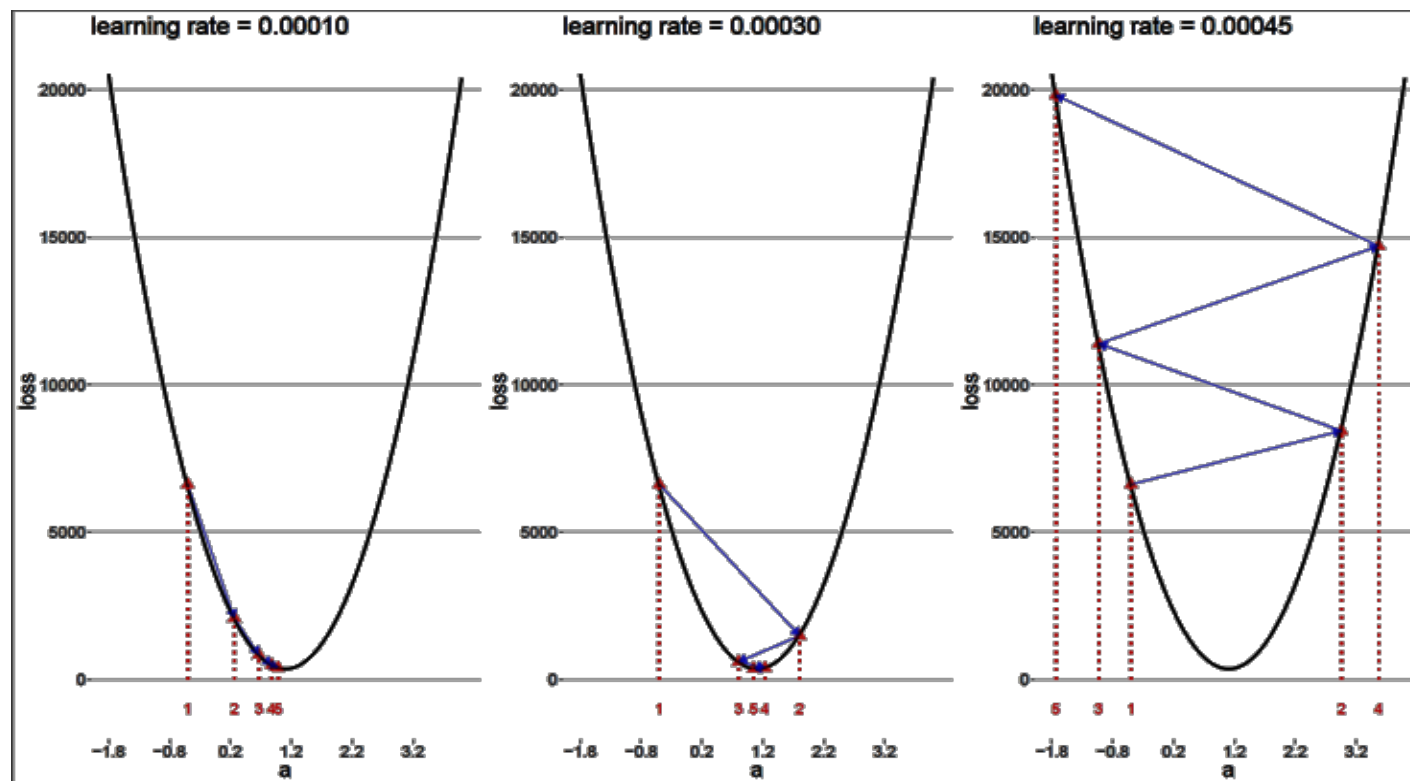
- gradients_weights /= N

- gradients_biases /= N

weights = weights - α * gradients_weights

biases = biases - α * gradients_biases

Learning Rate



Python Code

- Solving by hand
 - $X = [2, 4, 5, 6, 6, 8, 8, 9, 9]$
 - $y = [1, 4, 4, 5, 6, 7, 6, 8, 10]$
- Gradient Descent
- Stochastic Gradient Descent
- Convergence and Learning Rate

Inverse matrix technique

$$y = \beta X$$

$$\mathbf{x} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}$$

Inverse matrix technique

$$y = \beta X$$

$$\mathbf{x} = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}$$

$$\mathbf{e}(\beta) = \mathbf{y} - \mathbf{x}\beta$$

Inverse matrix technique

$$MSE(\beta) = \frac{1}{n} \sum_{i=1}^n e_i^2(\beta)$$

$$MSE(\beta) = \frac{1}{n} \mathbf{e}^T \mathbf{e}$$

$$MSE(\beta) = \frac{1}{n} \mathbf{e}^T \mathbf{e}$$

$$= \frac{1}{n} (\mathbf{y} - \mathbf{x}\beta)^T (\mathbf{y} - \mathbf{x}\beta)$$

$$= \frac{1}{n} (\mathbf{y}^T - \beta^T \mathbf{x}^T) (\mathbf{y} - \mathbf{x}\beta)$$

$$= \frac{1}{n} (\mathbf{y}^T \mathbf{y} - \mathbf{y}^T \mathbf{x}\beta - \beta^T \mathbf{x}^T \mathbf{y} + \beta^T \mathbf{x}^T \mathbf{x}\beta)$$


$$\begin{aligned} \nabla MSE(\beta) &= \frac{1}{n} (\nabla \mathbf{y}^T \mathbf{y} - 2\nabla \beta^T \mathbf{x}^T \mathbf{y} + \nabla \beta^T \mathbf{x}^T \mathbf{x}\beta) \\ &= \frac{1}{n} (0 - 2\mathbf{x}^T \mathbf{y} + 2\mathbf{x}^T \mathbf{x}\beta) \\ &= \frac{2}{n} (\mathbf{x}^T \mathbf{x}\beta - \mathbf{x}^T \mathbf{y}) \end{aligned}$$

$$\hat{\beta} = (\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{y}$$

Prediction intervals

$$E = t_{.025} * S_e * \sqrt{1 + \frac{1}{n} + \frac{n(x_0 + \bar{x})^2}{n(\sum x^2) - (\sum x)^2}}$$

$$S_e = \sqrt{\frac{\sum (y - \hat{y})^2}{n - 2}}$$





Python Code

- Analytical approach
- Prediction intervals



Non-linear Regression

- **Polynomial:** $y = \beta_0 + \beta_1 x + \beta_2 x^2 + \dots + \beta_n x^n$
- **Exponential:** $y = \beta_0 e^{\beta_1 x}$
- **Logarithmic:** $y = \beta_0 + \beta_1 \ln(x)$
- **Sigmoidal:** $y = \frac{L}{1+e^{-k(x-x_0)}}$ (used in logistic growth models)
- **Power Law:** $y = \beta_0 x^{\beta_1}$






Non-linear Regression



- Transform data
- Find optimal parameters
- Transform parameters





Non-linear Regression (Overfitting)



Overfitting




Bias and Variance Tradeoff




(a) Low bias, low variance.



(b) High bias, low variance.

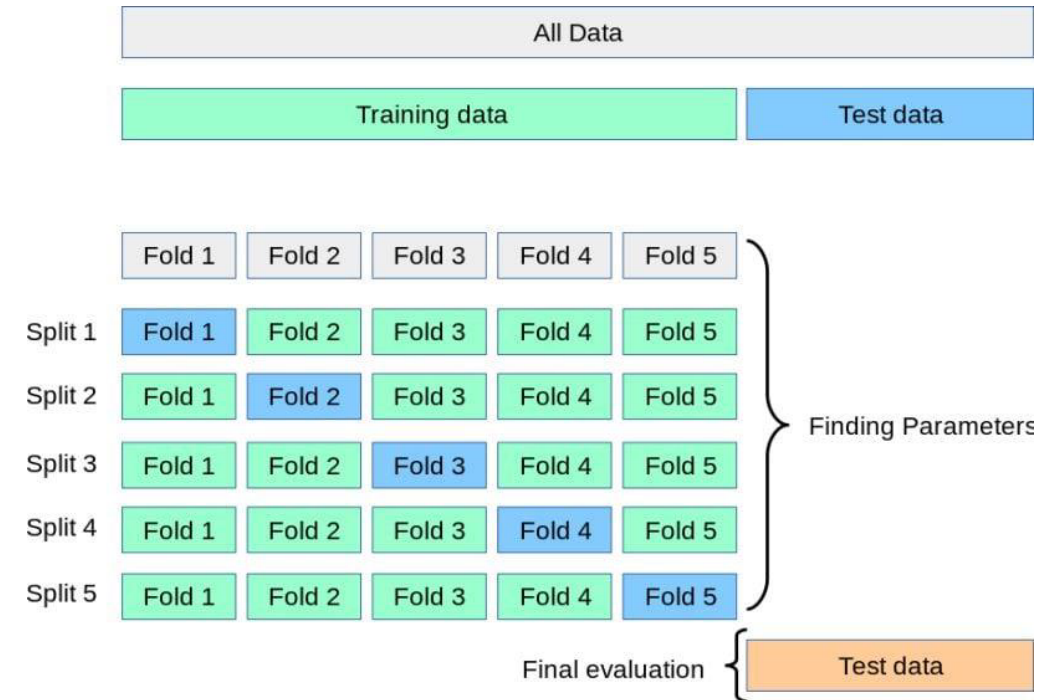
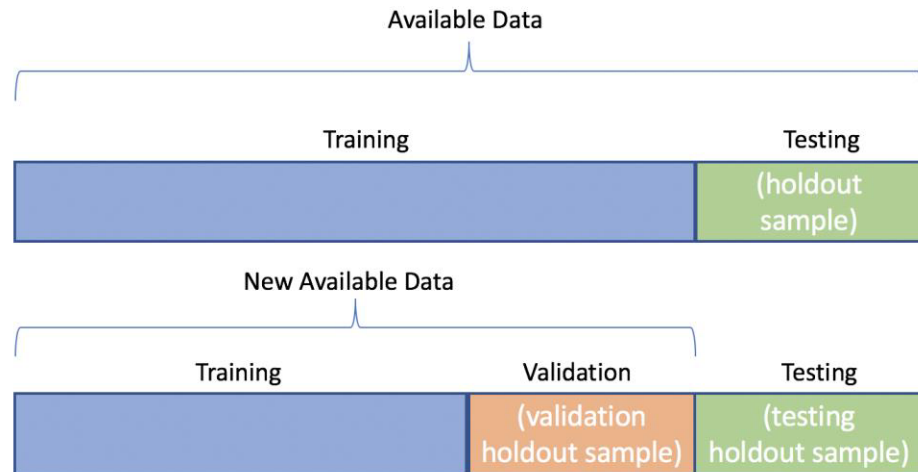


(c) Low bias, high variance.

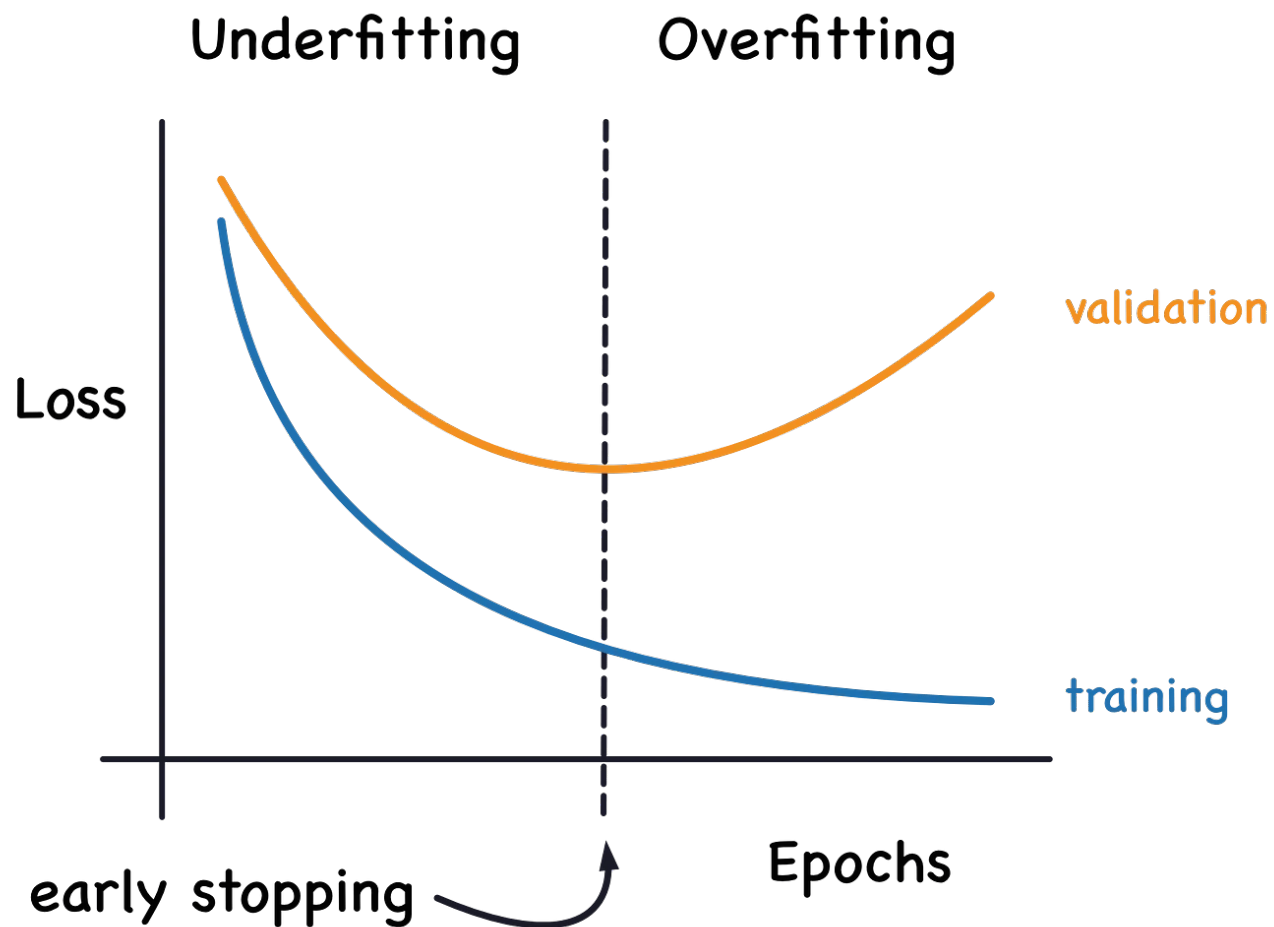


(d) High bias, high variance.

Train-Test Split-Cross Validation



Early Stopping



L2 Regularization

$$\text{Loss} = \frac{1}{2n} \sum_{i=1}^n (y_i - \hat{y}_i)^2 + \frac{\lambda}{2} \sum_{j=1}^p \beta_j^2$$

$$\frac{\partial \text{Loss}}{\partial \beta_j} = -\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i) x_{ij} + \lambda \beta_j$$

$$\beta_j = \beta_j - \eta \frac{\partial \text{Loss}}{\partial \beta_j}$$

Python

- Non-linear Regression
- L2 Regularization

Homework

- Download dataset from github (data-regression.csv)
- Split data into train and test.
- Train your model based-on one of common mathematical function.
- Use test data for evaluation your model.
- All code should write from scratch.